

Principles for modelling packages

TBD

2018-08-08

Contents

Chapter 1

Intro

Rule 1: Always spell it *modelling*, never *modeling*.

This document is targetted at R developers writing new packages for modelling.

For users interesting in learning more about using R for modelling in practice, we recommend X and Y

Chapter 2

Conceptual overview of modelling

- what is a model: models, estimands, estimators and model specifications
- what do we do with models
- how do fit models
- once we have a fit model, how do we predict or do inference
- the difference between working with a single fit vs a set of fits. LASSO example: wanting to use the coefficients for prediction vs wanting to see the order in which features enter the model

Chapter 3

Getting started on a modelling package

General dos:

- Export the `predict()` method
- Document the `predict()` method
- Use `match.arg()` for categorical arguments
- Validate the arguments to all your functions, especially our data

General dont's

-

Chapter 4

Model objects

- some explanation of why and how to save the function call
- generally what kinds of things should go into a model object, giving model objects a class so other people can extend them
- S3 object creation and validation for model building a la Advanced R
- Model classes beyond lists. when is S4 worth it? when is R6?
- Every modeling function should include its package version in its data object I will now save my models as a list of three objects: model, data, and `sessioninfo::session_info()`

Chapter 5

Data Specification

- formulas, model.frame, term objects, etc
- data / design matrix specification - **recipes**

habit: get the df right, then $y \sim .$ in the formula. would be nice to still see the features in the call?

- ask users to use data.frames and tibbles, not matrices.

5.1 Formulas

5.1.1 Testing formulas

https://github.com/alexpghayes/formulize/blob/master/tests/testthat/test_formula.R

minimum set of formula tests (based on mtcars example dataset):

- using `as.factor()` inline `mpg ~ as.factor(hp)`
- using `as.character()` inline `mpg ~ as.character(hp)`
- intercept only `mpg ~ 1`
- no intercept `mpg ~ disp + hp + drat - 1`
- implicit intercept `mpg ~ disp + hp + drat`
- explicit intercept `mpg ~ disp + hp + drat + 1`
- polynomials with 1 term `mpg ~ disp + hp + poly(drat, 1)`
- polynomials with multiple terms `mpg ~ disp + hp + poly(drat, 3)`
- natural splines with 1 term `mpg ~ disp + hp + ns(drat, 1)`
- natural splines with multiple terms `mpg ~ disp + hp + ns(drat, 3)`
- explicit interactions `mpg ~ drat + hp + drat:hp`
- dot `mpg ~ .`
- star `mpg ~ hp * drat`
- as.is `mpg ~ hp + I(drat^2)`
- multiple response `cbind`
- multiple responses as matrix `y ~ x` where `y` is a matrix
- multiple predictors as matrix `y ~ x` where `x` is a matrix
- multiple predictors and responses together `y ~ x`, both `x`, `y` matrices
- transformed response `log(mpg) ~ hp + drat`

optional / to: - `survival::Surv` and `survival::strata` objects

Chapter 6

Functional programming principles

calls to fit should be pure: i.e. no side effects like plotting, and especially no plotting with invisible object
return - side effects: useful in interactive mode, irritating in programmatic mode

- type safety, particularly of returned objects
- type safety with respect to single fits vs sets of fits

Chapter 7

Data

specification, exported data, and data sets used internally in a package

- using data from the package in tests
- using data from *other* packages in tests

Chapter 8

Documentation

- vignette should include not only the coefficients as output in an example, but also those coefficients written up as a general latex model and as a latex model with those specific coefficients substituted in
- **show** your example data in the README so users immediately see the structure

function to write out model form and fitted model in latex for sanity checking: some sort of `model_report` / `model_form` generic. think `report` generic or `write.model` may be coming to `fable/forecast` soon.

it's a bad idea to expect users to learn the *math* for your model from function level documentation, or math presented in ascii or unicode or poorly rendered latex.

show write out the math in a nicely formatted vignette, and then clearly describe the connection between code objects and math objects there as well

documenting arguments:

- **data**: super important to document acceptable **types** and formats, and highly recommend provided a dataset in `data/` with this format so the user can see exactly what they need to provide.

bad doc: The dataset to fit on the model on better doc: A data.frame or tibble with one row per observation and one column per features. For example, `mtcars` is in this format, but `messy_data` is not. It is okay to specify a matrix so long as it can coerced to a tibble. etc etc

Chapter 9

Testing

- testing against existing software - say a Matlab implementation
- saving long running models in `R/sysdata.rda` with `usethis::use_data(model_obj, internal = TRUE)`

Chapter 10

Workflow

10.1 Prediction

1. feature engineering
2. ML wizardry
3. more feature engineering
4. ???
5. predictions

10.2 Inference

1. Clean data
2. Specify model
3. Fit model
4. Check that model fitting process converged / worked
5. Check statistical assumptions of model

KEY part that always gets left out: working with multiple modellings

Chapter 11

Interface

- user friendly interfaces

good and bad existing idioms

- methods to implement
- examples of tried and true workflows

methods to implement - note on plotting: Should be easy to get the values plotted so others can make their own plots

TWO DISTINCT ISSUES THAT GET RESOLVED IN FORMULAE:

design matrix specification

model specification. (a la `fGarch::garchFit(~arma(1, 1) + garch(1, 0))`)

Chapter 12

Low and high level interfaces

- high level versus low level interface
- programmatic versus interactive use

when you should use which

examples: - high level: keras, brms - low level: tensorflow, stan

Chapter 13

Interactive modelling

what most people do different because there's a person looking at stuff as opposed to programmatic model when it's just code interacting with the model with no human involved

this is a chapter mostly to remind us to think of differences between the two and how they might be important in terms of interface

Chapter 14

Programmatic modelling

i.e. interacting with models programmatically

examples: - packages that export a model from someone to use a la `botornot` - models sitting behind a Plumber API - etc

Chapter 15

Vocabulary

useful functions that all modelling package developers should be aware of
`all.names(terms_object)` `all.vars(terms_object)`

15.1 model frame stuff

`model.frame` `mode`

15.2 na.action stuff

15.3 quoting operators

Chapter 16

Naming things

16.1 How to name function arguments

16.2 How to name model components

some standard names

currently lots of work happening in this realm in `broom` and the Stan community

<https://github.com/tidymodels/broom/issues/452>

Chapter 17

Danger Zone

little things to include somewhere: - the danger of misspecified arguments disappearing into ...

don't give your model `lm` class and count on other stuff to magically work. you should implement methods specifically for your class, and if that is just wrapping `lm` internals, great, but take explicit control

provide an example of extensive formula tests: include a log term in tests to catch formula edge cases

17.1 Warnings / places to use care

`model.frame()` is not the be all end all

double evaluation issue in

```
fit <- lm(hp ~ log(mpg), mtcars)
predict(fit, newdata = model.frame(fit))
```

```
## Error in eval(predvars, data, env): object 'mpg' not found
```

don't do `predict(object, newdata = model.frame(object))` since this will breaaaaaaaak

17.2 Anti-patterns

17.2.1 Using the default method of a generic

i.e. funneling everything into `confint.default`. Default methods for new generics should throw an error.

What do to about existing generics?

Key principle here: want to *guarantee* to the user that they are getting the right numbers

don't funnel everything through `augment_columns` and add special cases slowly - v hard to maintain. much better to have individualized S3 methods with *consistent behavior* abstracted out into small helpers.

related idea to the belong: have enough classes! using inheritance appropriately.

17.2.2 the documentation that isn't documentation and doesn't feature an actual use case

- *cough* ?MASS::predict.rlm *cough*
- misisng doc ?MASS::predict.polr()

17.2.3 Never use missing arguments

because people have to write more code to pass the objects they want

predict(m) and predict(m, newdata = NULL) should do the same thing you should test this

17.2.4 special casing everything through one workhorse function instead of using S3 methods

basically `augment_columns`

don't funnel everything through `augment_columns` and add special cases slowly - v hard to maintain. much better to have individualized S3 methods with *consistent behavior* abstracted out into small helpers.

17.3 Things to be aware of

If you call things `data` or `df` and users fail to specify data arguments, R will try and perform dataframe operations on the *functions* `data` and `df`. The resulting error messages for this can be cryptic. In this case you may wish to write an informative error message with a hint:

```
Error: Can't *do data frame thing* on a function.  
Are you sure you passed a tibble to *argument*?
```

Chapter 18

References

- bdr's model fitting functions in r